Ganesh Pawar

Mumbai, India | 8767880796 | ganeshppawar864@gmail.com | LinkedIn | GitHub | Portfolio

Profile

AI/ML Engineer specializing in **GenAI systems engineering** from classical ML pipelines to production-grade agentic AI. I design and ship **multi-agent LangGraph systems**, **RAG pipelines** (ChromaDB, FAISS, Qdrant), and **LLMOps-monitored FastAPI deployments** with LangSmith tracing, hallucination evaluation, and structured output validation. Hands-on across the full stack: **PyTorch/Scikit-learn** for modelling, **vector retrieval** for context, **LangChain/LangGraph** for orchestration, and **Docker + CI/CD** for production. I don't bridge the gap between research and deployment **I close it.**

Education

Bachelor of Engineering, Computer Science (AI & ML) — 2022 – 2026
Vishwaniketan's iMEET, Mumbai University

Experience

Sunitiq | Remote — Apr 2026 – Present
Machine Learning Intern

- **Engineered end-to-end predictive analytics workflows** using Python, PyTorch, and Scikit-learn — delivering production-ready regression and classification models for demand forecasting and business outcome prediction under agile practices.
- Built robust **feature engineering pipelines** (imputation, encoding, scaling, selection) with **cross-validation & hyperparameter tuning**, reducing manual analysis effort by **40%** and improving model generalization.
- **Automated model evaluation and experiment tracking** via structured logging, cutting iteration cycles by **30%** and enabling reproducible, audit-ready ML outputs.

InternHack | Remote — Jan 2026 – May 2026
Data Analytics Intern

- Built **ETL pipelines** in Python and SQL, **reducing manual reporting effort by 40%** and surfacing actionable anomalies via EDA integrated into downstream analytics workflows.
- Designed **Power BI dashboards** with DAX measures and KPI tracking, translating AI pipeline outputs into real-time operational insights for product and business stakeholders.

Technical Skills

Languages & DB	Python (Pandas, NumPy, Matplotlib, Seaborn, Plotly), SQL, PostgreSQL, MySQL
Generative AI & LLMs	LangChain, LangGraph, RAG Pipelines, Prompt Engineering, Function Calling, Structured Outputs, LLM Fine-tuning, LLM Evaluation (hallucination, quality), OpenAI API, Gemini API, Anthropic API, Mistral AI, Hugging Face
Embeddings & Vector DB	ChromaDB, FAISS, Qdrant, Semantic Search, BGE Embeddings, Embedding Pipelines, NER, Text Classification
MLOps & LLMOps	LangSmith (tracing, eval), Experiment Tracking, Model Monitoring, Docker, Git, GitHub CI/CD
ML & Data Science	Scikit-learn, PyTorch, TensorFlow/Keras, CNN, Feature Engineering, Regression, Classification, Clustering, Hyperparameter Tuning, A/B Testing
Backend & Platforms	FastAPI, Streamlit, Celery, Redis, Playwright, DuckDB, BigQuery (GCP), Power BI (DAX), n8n, Jupyter

Projects

VoiceOps - AI-Powered Voice Software Installation Agent | GitHub

- Designed an **8-agent LangGraph Plan-and-Execute pipeline** (Speech, Intent, Planner, Browser, Download, Install, Monitor, Notify) converting natural language voice commands into **fully autonomous software installations** end-to-end.
- Integrated **faster-whisper STT, Mistral AI, Qdrant RAG, Playwright browser automation**, SHA-256 verification, Celery + Redis, PostgreSQL, and SSE/WebSocket notifications, deployed via FastAPI + Docker with CI/CD.

Agentic Data Analyst - Autonomous Business Intelligence System | GitHub

- Architected an **8-agent LangGraph multi-agent system** (Planner, Analyst, Query, Visualization, Forecasting, Reporting) with RAG (ChromaDB), Redis memory, NL-to-SQL, and function calling, autonomously detecting anomalies and generating KPI reports for non-technical stakeholders.
- Shipped a **full LLMOps-monitored production stack** on FastAPI, PostgreSQL, DuckDB, Prophet, ARIMA, Docker, and LangSmith, with hallucination scoring, quality eval, time-series forecasting, and PDF/HTML executive reporting.

HealthVerse AI - Intelligent Medical Report Analysis System | GitHub

- Architected a **RAG-powered GenAI system** fusing CNN deep learning with LLMs for medical report understanding, processing **5,000+ records** and reducing retrieval latency by **60%** via optimized vector search (ChromaDB + BGE Embeddings).
- Achieved **93.6% diagnostic classification accuracy** and deployed end-to-end via FastAPI + Streamlit, covering the full AI lifecycle from data ingestion through deployment and **LLM output monitoring**.

Certifications & Achievements

- **Data Science, ML, Deep Learning & NLP** — Udemy | *Krish Naik* — statistical ML, EDA, supervised/unsupervised learning, deep learning architectures, and NLP with production-oriented project implementation.
- Shipped **3 production-grade GenAI systems** multi-agent BI automation, RAG-powered LLM orchestration, and a voice-controlled agentic AI each with LLMOps monitoring, hallucination evaluation, and stakeholder-ready reporting.
- Achieved **93.6% diagnostic accuracy** on real-world clinical data; **40% reduction** in manual reporting effort and **60% reduction** in retrieval latency across deployed AI systems.